

Shopware AI Agent

Autonomous Plugin Development with LLMs

SCUC 2026 • @jankal • nuonic Digital

Why?

-  Plugin development is **repetitive**
 - boilerplate: `composer.json`, Plugin class, `services.xml` ...
 - same patterns over and over
-  Context switching is **costly**
 - docs, Admin API, CLI, code, browser
-  LLMs are great at code generation
 - but they lack **execution context**

The Gap


An LLM that can only *generate* code is a **fancy autocomplete**.

An LLM that can **write**, **execute**, **validate**, and **see** the result is an **autonomous agent**.

→ We gave Claude a Shopware instance to play with.



Architecture: Sidecar Pattern



Why a Sidecar?

- Shared Volume = **zero-latency** file I/O
 - Agent writes PHP → instantly visible to Shopware
- `ReadWriteOnce` PVC is sufficient (same Pod)
 - no expensive RWX network storage needed
- **K8s Exec API** for command execution
 - no SSH, no network hops
- Independent lifecycle
 - Shopware crashes? Agent still runs.

Tech Stack

Component	Technology
Agent Runtime	Nuxt 4 + Bun
LLM	Claude Opus 4 (Anthropic SDK)
Tool Bridge	MCP (Model Context Protocol)
K8s Integration	<code>@kubernetes/client-node</code>
Conversations	PostgreSQL + Drizzle ORM
Validation	<code>shopware-cli</code> (PHPStan, CS)
Screenshots	Playwright (Chromium)
Frontend	Nuxt UI + Tailwind CSS

 **How?**

Native Tools

The agent has **6 built-in tools** for direct Shopware interaction:

Tool	What it does
<code>write_file</code>	Write to shared volume (with diff)
<code>read_file</code>	Read from shared volume
<code>list_directory</code>	Browse the filesystem
<code>search_files</code>	Glob + content search
<code>exec_command</code>	Shell commands via K8s Exec API
<code>validate_extension</code>	PHPStan, code style via <code>shopware-cli</code>

K8s Exec API

```
// k8s-exec.ts (simplified)
export async function execInContainer(command: string) {
  const kc = new k8s.KubeConfig()
  kc.loadFromCluster()

  const exec = new k8s.Exec(kc)

  await exec.exec(
    NAMESPACE,           // from Downward API
    POD_NAME,             // from Downward API
    'shopware-app',       // target container
    ['/bin/sh', '-c',
     `cd /var/www/html && su -s /bin/sh www-data -c '${command}'`],
    stdout, stderr,
    null,                 // no stdin
    false,                // no tty
    (status) => resolve(status)
  )
}
```

No SSH. No network hops. Just the **Kubernetes API**.

MCP Server Tools

Three **MCP servers** extend the agent's capabilities:

admin-mcp ([@shopware-ag/admin-mcp](#))

→ Products, categories, orders, themes, media via Admin API

docs-mcp (custom, Algolia)


→ Search developer.shopware.com documentation

screenshot-mcp (custom, Playwright)

→ Take screenshots, click elements, fill forms


MCP = Model Context Protocol — open standard for LLM tool integration

MCP Bridge



Tools are **discovered dynamically** — the bridge merges all MCP tools into the LLM tool palette at runtime.

The Agentic Loop



Up to **25 rounds** per conversation turn.
SSE streaming for real-time UI updates.

Tool Tracking & Revert

Every `write_file` is tracked with full diffs:

```
// dispatchTool (simplified)
case 'write_file': {
  const result = await writeFileOp(input.path, input.content)
  // result.diff → unified diff
  // result.previousContent → for revert

  await db.insert(toolActions).values({
    conversationId,
    toolName: 'write_file',
    diff: result.diff,
    previousContent: result.previousContent,
    isReversible: true,
  })
}
```

→ Users can **revert any file write** from the UI.

System Prompt

The agent has a detailed **German system prompt** with:

-  Shopware plugin directory structure
-  Required `composer.json` fields
-  CLI command reference (`plugin:refresh`, `cache:clear`, ...)
-  Best practices (`strict_types`, PSR-4, namespaces)
-  Admin API tool documentation
-  Validation workflow (validate → fix → repeat)
-  Screenshot workflow (compile → screenshot → show)



"Create a plugin that adds a custom CMS element"

What Happens Behind the Scenes

1.  User sends prompt via SSE stream
2.  Claude plans the plugin structure
3.  `write_file` × 5–8 calls
composer.json, Plugin.php, services.xml, CmsElement, ...
4.  `exec_command: plugin:refresh && plugin:install`
5.  `exec_command: plugin:activate && cache:clear`
6.  `validate_extension`: PHPStan + code style
7.  Auto-fix any validation errors
8.  `screenshot`: verify result in storefront
9.  Claude reports back with screenshot

Embedded in Shopware Admin

The agent UI runs as a **Shopware App** (not a plugin):

- `manifest.xml` with `<admin><module>`
- **Zero PHP code** — purely declarative
- iframe in Shopware Admin under Settings
- `postMessage( 'sw-app-loaded' )` handshake
- Works even if Shopware is broken
(agent runs in a separate container)

Security

Filesystem

- Path traversal prevention (`safePath()`)
- All paths relative to `/var/www/html`

Kubernetes

- Namespace-scoped RBAC (ServiceAccount)
- Only `pods/exec` permission in own namespace
- Resource limits (500m CPU, 512Mi RAM)

Agent

Lessons Learned

-  **Context is everything**
Tools > more tokens. Let the LLM *do*, not just *say*.
-  **Validation loops are essential**
LLMs make mistakes. Automated validation catches them.
-  **Visual feedback closes the loop**
Screenshots let the LLM self-correct UI issues.
-  **MCP is the right abstraction**
Swappable tool servers, standardized protocol.
-  **Sidecar > microservice**
Shared volumes beat API calls for file I/O.

What's Next?

-  Multi-tenant SaaS (scale-to-zero with KEDA)
-  Auth between Shopware  Agent
-  Git export to production repos
-  Automated test generation
-  More MCP servers (GitHub, CI/CD, monitoring)

Questions?

@jankal • nuonic Digital • SCUC 2026